

Sberbank Russian Housing Market - Top 1% Solution

**Team
Chase Edge, Grant Webb
And Brandy Freitas**



Project Overview

Project Overview

- Kaggle Competition
- Predict housing prices in Moscow during July 2015 to May 2016 using data from August 2011 to June 2015
- Data includes housing transaction information (e.g. square meter, number of rooms and build year), neighborhood details and macroeconomic information

Machine Learning Checklist

1. Frame the Problem and Look at the Big Picture
2. Get The Data
3. Explore the Data
4. Prepare the Data
5. Short List Promising Models
6. Fine-Tune the System

Look at the Big Picture

Value to Sberbank

- Mitigate Risk
 - Help avoid overlending (issuing mortgages in excess of the value of the home)
- Valuation
 - Help banks value their portfolio
- Predictions
 - To give confidence to renters, developers and lenders when they sign a lease or purchase a building

Data Exploration

Data Overview

Housing Data

- Number of Observations
 - Training - 30,471
 - Test - 7,662
- Number of Features - 290

Macro Data

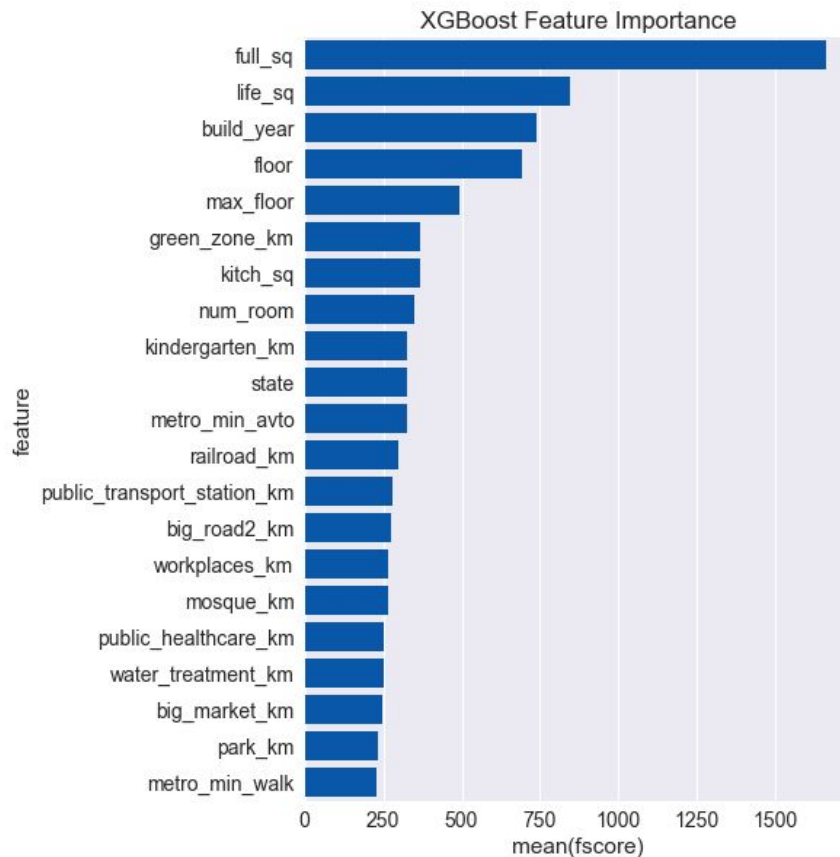
- Number of Observations - 2,484
- Number of Features - 100

Data Overview - Most Important Features

- Initial Thoughts
 - Square Meters
 - Number of Rooms
 - Build Year
 - Location

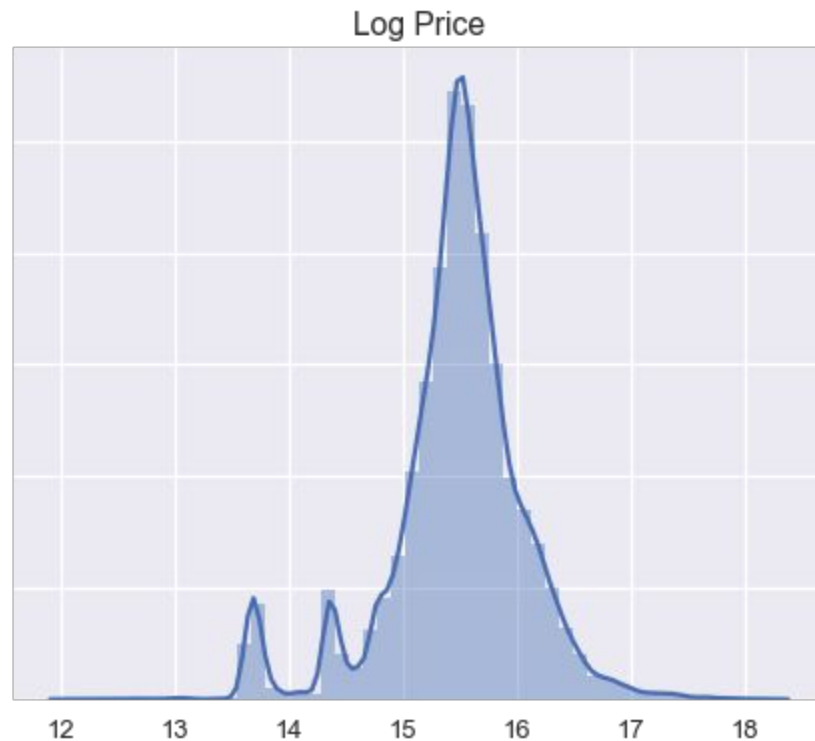
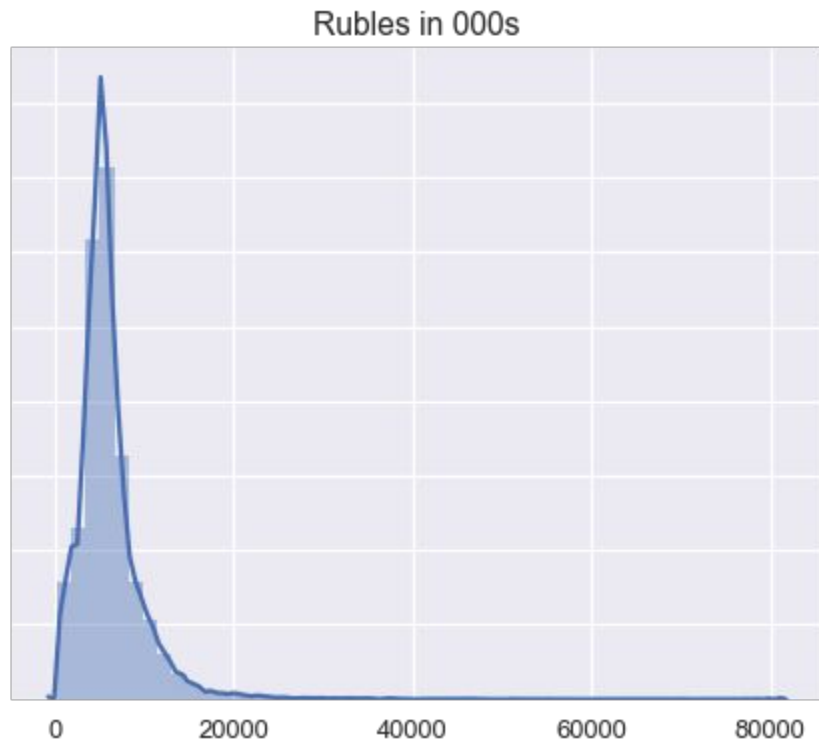
Data Overview - Most Important Features

- Initial Thoughts
 - Square Meters
 - Number of Rooms
 - Build Year
 - Location



Housing Price


Distribution of Prices




Housing Price

Distribution of Price / Meter Sq


Rubles in 000s



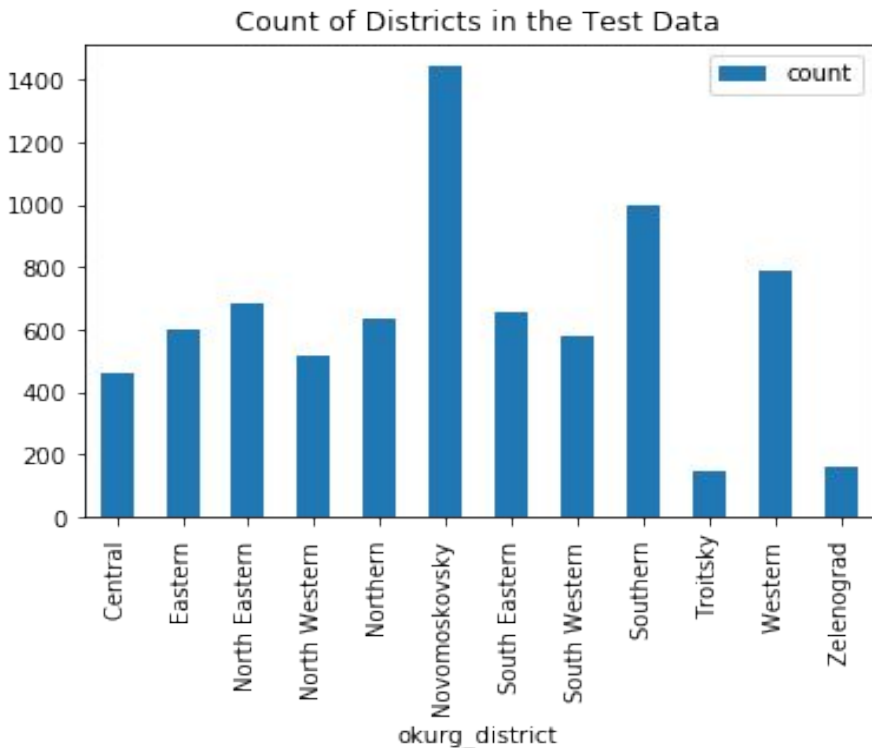
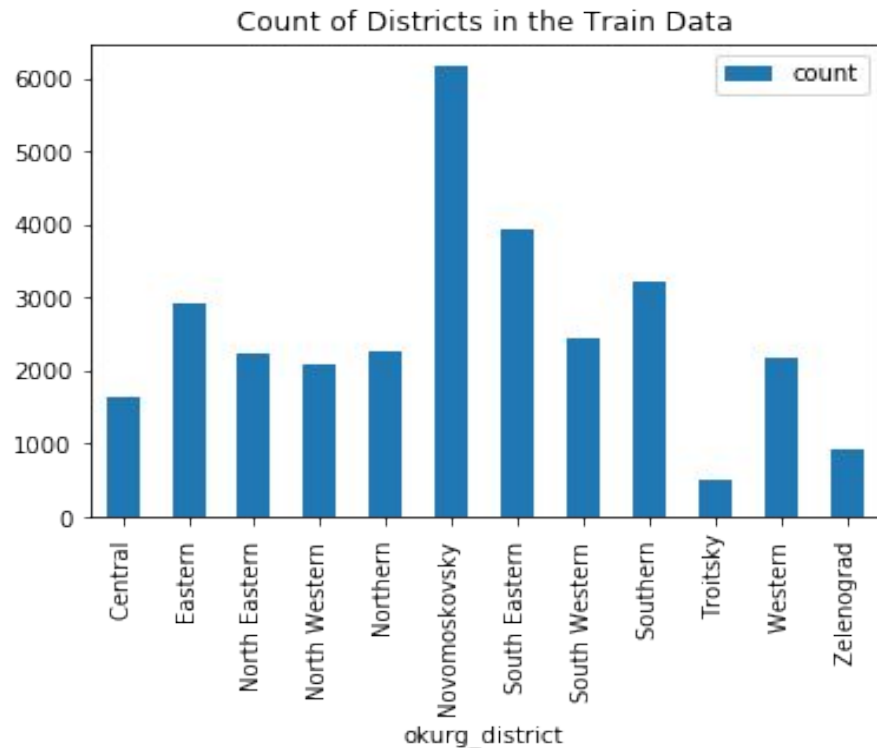
Log Price



Prices by District




Observations by District




Data Preparation

Missingness in Top Features

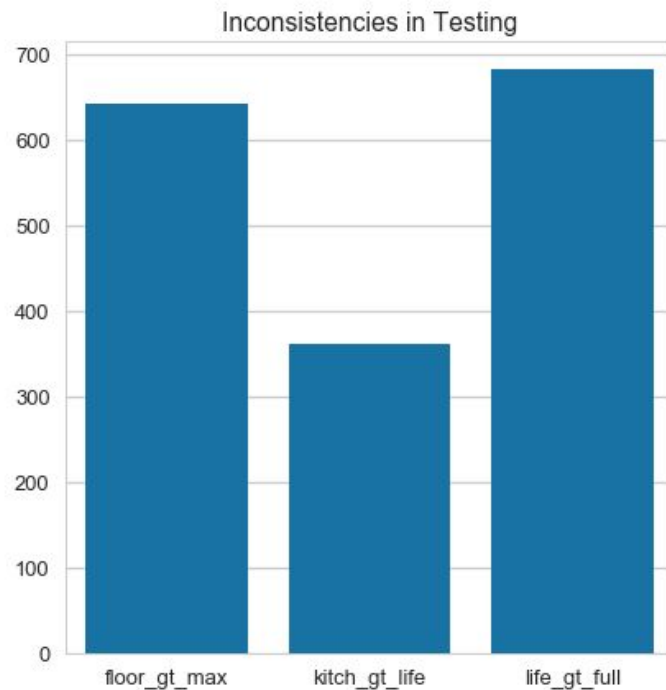
Percentage of Missingness in Training





Percentage of Missingness in Testing




Data Inconsistencies



Data Inconsistencies



Imputations - KNN by Neighborhood



kindergarten_km
park_km
public_transport_station_km
public_transport_station_min_walk
water_km
mkad_km
ttk_km
sadovoe_km
bulvar_ring_km
kremlin_km
big_road1_km
big_road2_km
railroad_km
oil_chemistry_km
nuclear_reactor_km
radiation_km
power_transmission_line_km
thermal_power_plant_km
ts_km
basketball_km
hospice_morgue_km
big_church_km
church_synagogue_km
mosque_km
museum_km
exhibition_km
catering_km

11852
11852
11851
11852
11851
11852
11852
11852
11852
11852
11852
11852
11852
11852
11852
11852
11852
11852
11849
11852
11852
11852
11852
11852
11852
11852
11852

| | | count |
|------------|-----------|-------|
| kremlin_km | | |
| | 20.549464 | 976 |
| | 0.072897 | 603 |
| | 23.373697 | 582 |
| | 20.666814 | 364 |
| | 22.222434 | 319 |
| | 29.133765 | 288 |
| | 25.735256 | 282 |
| | 15.869044 | 275 |
| | 21.609733 | 254 |
| | 18.752843 | 232 |
| | 25.595974 | 229 |
| | 19.763938 | 215 |
| | 22.567655 | 202 |

Feature Engineering & Insights

Feature Engineering

$Sq_diff = Full_sq - Kitch_sq$


$Floor_ratio = Floor / Max_floor$

$Life_to_full = Life_sq / Full_sq$


$Kitch_to_full = Kitch_sq / Full_sq$

$Room_size = Life_sq / Num_room$

Month = month of sale



Decline in Russian Housing Prices



Dealing with the Decline

- **Problem**

- Housing prices declined in 2015 and 2016
- Model predicts values that are too high

- **Solutions**

- Incorporate economic data
- Make downward adjustments to predicted values
- *Adjust prices for fluctuations in the market based on a price index*

Price Index

- The Data
 - Russian government statistics on the monthly rental prices within Moscow
- The Index
 - 3 month rolling average of 3 bed, 2 bed and 1 bed rentals in Moscow
 - Indexed to the start of the training data (August 2011)
 - Averaged the 3 indices
- The Application
 - Adjust all prices in the training data for changes in the index (divide by index)
 - Model predicts prices as if they occurred in August 2011
 - Adjust the predicted values for the index (multiply by index)

Price Index



Price Index - Example

- Training Transaction

- Date - April 2013
- Price - RUB 5,693,972
- Index - 1.13
- Adjusted Price - RUB 5,038,913

- Predicted Value

- Date - April 2016
- Price - RUB 3,902,007
- Index - 1.06
- Adjusted Price - RUB 4,136,127

Short-list Promising Models

KISS - OLS with 2 Variables


Features = full_sq, sub_area

Response = log_price

OLS Regression Results

```
=====
Dep. Variable:          log_price    R-squared:                0.574
Model:                  OLS          Adj. R-squared:           0.572
Method:                 Least Squares    F-statistic:             270.5
Date:                  Tue, 30 May 2017    Prob (F-statistic):       0.00
Time:                  18:13:54          Log-Likelihood:          -9345.5
No. Observations:      29096            AIC:                    1.898e+04
Df Residuals:          28951            BIC:                    2.018e+04
Df Model:               144
Covariance Type:       nonrobust
=====
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-------------------------|---------|----------|---------|-------|--------|--------|
| Intercept | 14.8414 | 0.022 | 682.397 | 0.000 | 14.799 | 14.884 |
| full_sq | 0.0146 | 9.91e-05 | 147.232 | 0.000 | 0.014 | 0.015 |
| sub_area_Ajeroport | 0.1418 | 0.037 | 3.827 | 0.000 | 0.069 | 0.214 |
| sub_area_Akademicheskoe | 0.2088 | 0.032 | 6.483 | 0.000 | 0.146 | 0.272 |
| sub_area_Alekseevskoe | 0.1792 | 0.040 | 4.427 | 0.000 | 0.100 | 0.258 |



Your submission scored 0.38249, which is not an improvement of your best score. Keep trying!

KISS - OLS with 7 Variables

Features = full_sq, floor, max_floor, life_to_full, kitch_to_life, product_type, sub_area
Response = log_price

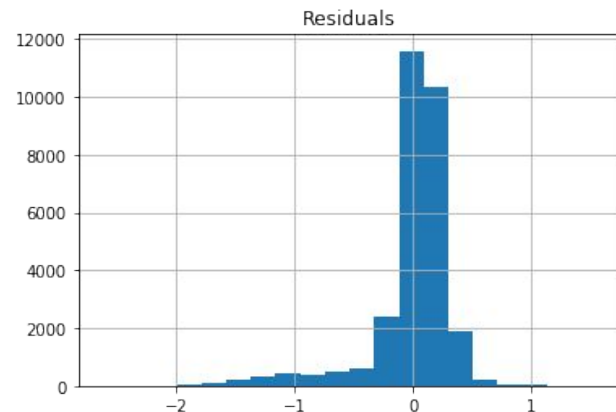
OLS Regression Results

```
=====
Dep. Variable:          log_price    R-squared:                0.577
Model:                  OLS          Adj. R-squared:           0.575
Method:                 Least Squares    F-statistic:             265.3
Date:                  Tue, 30 May 2017    Prob (F-statistic):       0.00
Time:                  23:59:56          Log-Likelihood:          -9220.6
No. Observations:      29096            AIC:                    1.874e+04
Df Residuals:          28946            BIC:                    1.998e+04
Df Model:              149
Covariance Type:       nonrobust
=====
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|----------------------------|---------|---------|---------|-------|---------|--------|
| Intercept | 9.9379 | 0.021 | 462.806 | 0.000 | 9.896 | 9.980 |
| full_sq | 0.0142 | 0.000 | 133.472 | 0.000 | 0.014 | 0.014 |
| floor | 0.0013 | 0.000 | 2.687 | 0.007 | 0.000 | 0.002 |
| max_floor | 0.0050 | 0.000 | 10.345 | 0.000 | 0.004 | 0.006 |
| life_to_full | -0.1116 | 0.024 | -4.733 | 0.000 | -0.158 | -0.065 |
| kitch_to_life | -0.1586 | 0.028 | -5.650 | 0.000 | -0.214 | -0.104 |
| product_type_Investment | 4.9614 | 0.009 | 553.891 | 0.000 | 4.944 | 4.979 |
| product_type_OwnerOccupier | 4.9764 | 0.014 | 349.646 | 0.000 | 4.949 | 5.004 |

simple_linear_052716_log_multiples.csv

0.34644



Tree-Based Models


Decision Tree

Unrestrained Decision Trees are known to overfit as shown:

Accuracy on training set 1.000

Accuracy on test set 0.224

The severe overfitting of the unrestrained Decision Tree!!



-- Grid Parameter Search via 5-fold CV

GridSearchCV took 106.79 seconds for 144 candidate parameter settings.

Model with rank: 1

Mean validation score: 0.537 (std: 0.006)

Parameters: {'min_samples_split': 2, 'max_leaf_nodes': None, 'max_depth': 10, 'min_samples_leaf': 10}

Model with rank: 2

Mean validation score: 0.537 (std: 0.006)

Parameters: {'min_samples_split': 20, 'max_leaf_nodes': None, 'max_depth': 10, 'min_samples_leaf': 10}

Model with rank: 3

Mean validation score: 0.537 (std: 0.006)

Parameters: {'min_samples_split': 10, 'max_leaf_nodes': 10, 'min_samples_leaf': 10}

Accuracy on training set 0.628

Accuracy on test set 0.568

No longer overfitting of the Decision Tree!!

```
# set of parameters to test
```

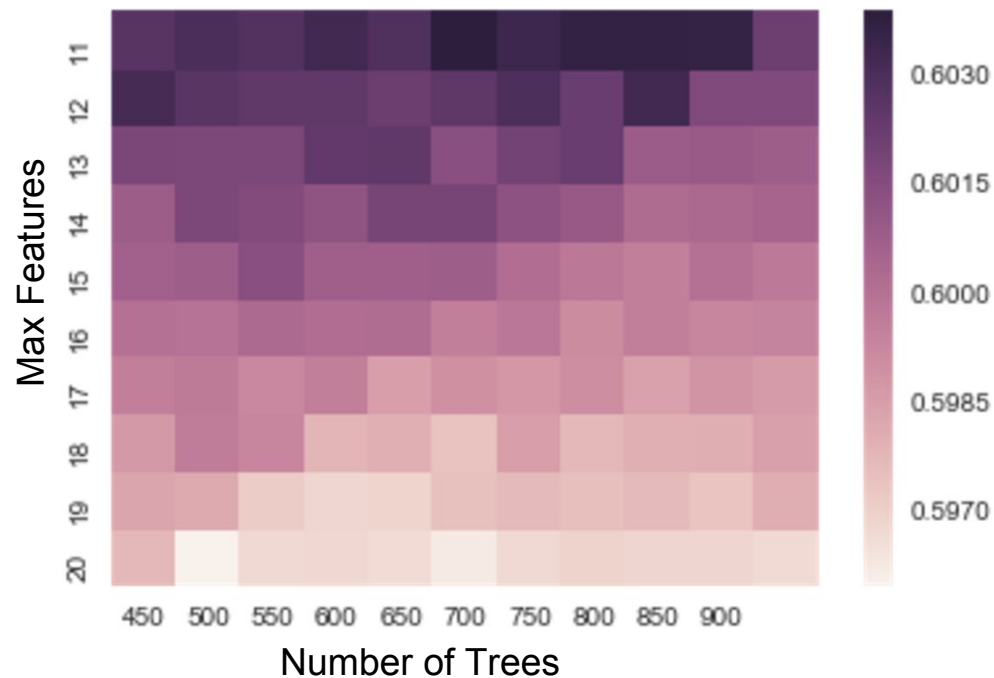
```
param_grid = {"min_samples_split": [2, 10, 20],  
              "max_depth": [None, 2, 5, 10],  
              "min_samples_leaf": [1, 5, 10],  
              "max_leaf_nodes": [None, 5, 10, 20],  
              }
```

Random Forest

Accuracy on training set 0.922

Accuracy on test set 0.581

Also overfitting of the unrestrained Random Forest!!



Accuracy on training set 0.947

Accuracy on test set 0.626

No longer overfitting of the Random Forest, as much!!

8 hrs to run with 3
cores processing!

XGBoost

- Log(Scaled Price / Meter Sq) as Target
- Fast - Training time 1 min 43 sec
- Best Results
- Less Interpretable

```
xgb_params = {  
    'eta': 0.01,  
    'max_depth': 5,  
    'subsample': 0.7,  
    'colsample_bytree': 0.3,  
    'objective': 'reg:linear',  
    'eval_metric': 'rmse',  
    'silent': 1,  
    'n_jobs': -1  
}  
  
xgb.train(xgb_params,  
          dtrain,  
          num_boost_round=1000,  
          evals=[(dval, 'val')],  
          early_stopping_rounds=20, verbose_eval=20)
```

| | | | | | | |
|----|-------|-------------|--|---------|-----|-----|
| 14 | ▲6 | RunningWolf |  | 0.31193 | 124 | 1mo |
| 15 | ▲901 | Chase Edge |  | 0.31212 | 35 | 2mo |
| 16 | ▲1092 | yrtchn |  | 0.31227 | 125 | 1mo |

Future Work

Next Steps

- Model Ensembling
- Feature Engineering
- Time Series Analysis on Pricing Index
- Cluster Analysis on Neighborhoods

Questions